

```
In [1]: import pandas as pd
df = pd.read_csv('titanic.csv')
df
```

```
Out[1]:
```

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Ci
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	100	0	A/5 21171	7.2500	
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	0	PC 17599	71.2833	
2	3	1	3	Heikkinen, Miss. Laina	female	26.0	0	0	STON/O2. 3101282	7.9250	
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	100	0	113803	53.1000	C
4	5	0	3	Allen, Mr. William Henry	male	35.0	0	0	373450	8.0500	
...
886	887	0	2	Montvila, Rev. Juozas	male	27.0	0	0	211536	13.0000	
887	888	1	1	Graham, Miss. Margaret Edith	female	19.0	0	0	112053	30.0000	
888	889	0	3	Johnston, Miss. Catherine Helen "Carrie"	female	NaN	1	2	W./C. 6607	23.4500	
889	890	1	1	Behr, Mr. Karl Howell	male	26.0	0	0	111369	30.0000	C
890	891	0	3	Dooley, Mr. Patrick	male	32.0	0	0	370376	7.7500	

891 rows × 12 columns



```
In [2]: df.isna().sum()
```

```
Out[2]: PassengerId      0
        Survived         0
        Pclass          0
        Name            0
        Sex             0
        Age            177
        SibSp           0
        Parch           0
        Ticket          0
        Fare            0
        Cabin          687
        Embarked        2
        dtype: int64
```

```
In [3]: mean_age = round(df['Age'].mean())
        mean_age
```

```
Out[3]: 30
```


```
In [4]: df['Age'].fillna(mean_age, inplace=True)
```

```
In [7]: from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()
df['Cabin'] = le.fit_transform(df['Cabin'])
df['Embarked'] = le.fit_transform(df['Embarked'])
df['Ticket'] = le.fit_transform(df['Ticket'])
df['Sex'] = le.fit_transform(df['Sex'])
df['Name'] = le.fit_transform(df['Name'])
df
```

```
Out[7]:
```

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Emb
0	1	0	3	108	1	22.0	100	0	523	7.2500	147	
1	2	1	1	190	0	38.0	1	0	596	71.2833	81	
2	3	1	3	353	0	26.0	0	0	669	7.9250	147	
3	4	1	1	272	0	35.0	100	0	49	53.1000	55	
4	5	0	3	15	1	35.0	0	0	472	8.0500	147	
...
886	887	0	2	548	1	27.0	0	0	101	13.0000	147	
887	888	1	1	303	0	19.0	0	0	14	30.0000	30	
888	889	0	3	413	0	30.0	1	2	675	23.4500	147	
889	890	1	1	81	1	26.0	0	0	8	30.0000	60	
890	891	0	3	220	1	32.0	0	0	466	7.7500	147	

891 rows × 12 columns



```
In [8]: mode_cabin = df['Cabin'].mode()[0]
mode_cabin
```

Out[8]: 147

```
In [9]: df['Cabin'].fillna(mode_cabin, inplace=True)
```

```
In [10]: mode_embarked = df['Embarked'].mode()[0]
mode_embarked
```

Out[10]: 2


```
In [12]: df['Embarked'].fillna(mode_embarked, inplace=True)
```

```
In [13]: df.isna().sum()
```

```
Out[13]: PassengerId    0
Survived              0
Pclass               0
Name                 0
Sex                  0
Age                  0
SibSp                0
Parch                0
Ticket              0
Fare                 0
Cabin                0
Embarked             0
dtype: int64
```

```
In [17]: import matplotlib.pyplot as plt
%matplotlib inline
plt.scatter(df['SibSp'], df['Survived'])
```

```
Out[17]: <matplotlib.collections.PathCollection at 0x21309780e80>
```



```
In [19]: mode_sibsp = df['SibSp'].mode()[0]
for x in range(0,891):
    if df['SibSp'][x] >= 20:
        df['SibSp'][x] = mode_sibsp
```

C:\Users\USER\AppData\Local\Temp\ipykernel_10164\4235250280.py:4: SettingWithCopyWarning:


A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
df['SibSp'][x] = mode_sibsp
```


```
In [20]: import matplotlib.pyplot as plt
%matplotlib inline
plt.scatter(df['SibSp'], df['Survived'])
```

Out[20]: <matplotlib.collections.PathCollection at 0x21309acd490>



```
In [22]: import matplotlib.pyplot as plt
%matplotlib inline
plt.scatter(df['Fare'], df['Survived'])
```

Out[22]: <matplotlib.collections.PathCollection at 0x21309b8ae20>



```
In [23]: mode_fare = df['Fare'].mode()[0]
for x in range(0,891):
    if df['Fare'][x] >= 400:
        df['Fare'][x] = mode_fare
```

C:\Users\USER\AppData\Local\Temp\ipykernel_10164\690323287.py:4: SettingWithCopyWarning:


A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
df['Fare'][x] = mode_fare
```

```
In [24]: import matplotlib.pyplot as plt
%matplotlib inline
plt.scatter(df['Fare'], df['Survived'])
```

Out[24]: <matplotlib.collections.PathCollection at 0x21309c0a220>



```
In [25]: x = df.drop(['Survived'], axis=1)
y = df['Survived']
```

```
In [26]: from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(x,y,test_size=0.2,random_stat
```

```
In [27]: from sklearn.preprocessing import MinMaxScaler
min_max_scaler = MinMaxScaler()
```

```
In [28]: x_train = min_max_scaler.fit_transform(x_train)
x_test = min_max_scaler.fit_transform(x_test)
```

```
In [29]: from sklearn.linear_model import LogisticRegression
model = LogisticRegression()
```

```
In [30]: model.fit(x_train,y_train)
```

```
Out[30]: LogisticRegression()
```

```
In [31]: y_pred = model.predict(x_test)
y_pred
```

```
Out[31]: array([0, 0, 0, 1, 1, 1, 1, 0, 1, 1, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0,
                1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0,
                1, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 1, 0, 1, 1, 1, 0, 1, 1, 0, 0, 1,
                0, 0, 0, 1, 1, 1, 1, 1, 0, 0, 1, 1, 1, 0, 0, 1, 1, 0, 0, 0, 1, 1,
                0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 1,
                1, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 1, 1, 0, 0, 1, 1, 1, 1, 0, 1, 0,
                0, 1, 0, 1, 1, 0, 0, 1, 0, 1, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0,
                0, 0, 0, 1, 1, 1, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 1, 0, 1, 0, 0,
                0, 1, 0], dtype=int64)
```

```
In [33]: from sklearn.metrics import classification_report
classification_report(y_test, y_pred)
```

```
Out[33]: '          precision    recall  f1-score   support\n\n      81          0.90          0.85         105\n      74\n\n      accuracy          0.82         179\n\n      0.82          0.80          0.80         179\n\n\n      weighted avg          0.82          0.82          0.81\n      179\n'
```